




# Introduction to Reinforcement Learning

Alexandre Carbonneau, M. Sc.



# Motivation






## Supervised learning example



## Unsupervised learning example



# RL examples



## Breakout

`https://www.youtube.com/watch?v=TmPfTpjtdgg`

# Introduction to Reinforcement Learning

Goal of this presentation:

- 1) Present a general introduction of reinforcement learning (RL).
- 2) Show one method to solve the RL problem: Q-learning (will implement in Python).

# RL vs machine learning

How is RL different than the typical machine learning framework.

- ▶ **No supervisor:** only a reward signal that can be sparse (i.e. delayed reward, not instantaneous).



# RL vs machine learning

How is RL different than the typical machine learning framework.

- ▶ **No supervisor**: only a reward signal that can be sparse (i.e. delayed reward, not instantaneous).
- ▶ Sequential **non i.i.d.** data (time is important).

# RL vs machine learning

How is RL different than the typical machine learning framework.

- ▶ **No supervisor**: only a reward signal that can be sparse (i.e. delayed reward, not instantaneous).
- ▶ Sequential **non i.i.d.** data (time is important).
- ▶ Agent's actions affect the subsequent data it receives.
  - Chess, self-driving cars, financial trading agent with a limit order book (market impact), etc...

# Rewards

Every step, the agent gets a scalar feedback signal  $R_t \in \mathbb{R}$  called the **reward** at step  $t$ .

- Agent's goal: select a sequence of actions to maximize **the cumulative future reward**.
- Actions may have long-term consequences.
  - ▶ Might be better to sacrifice immediate reward to gain more long-term rewards!

## Examples

- ▶ Shortest path problems  
Make  $R_t = -1$  at each step

## Examples

- ▶ Shortest path problems  
Make  $R_t = -1$  at each step
- ▶ Financial trading agent  
Make  $R_t$  the P&L from time  $t - 1$  to  $t$ .

## Examples

- ▶ Shortest path problems  
Make  $R_t = -1$  at each step
- ▶ Financial trading agent  
Make  $R_t$  the P&L from time  $t - 1$  to  $t$ .
- ▶ Game of chess  
 $\pm 1$  for winning/losing a game, else zero.

## Examples

- ▶ Shortest path problems  
Make  $R_t = -1$  at each step
- ▶ Financial trading agent  
Make  $R_t$  the P&L from time  $t - 1$  to  $t$ .
- ▶ Game of chess  
 $\pm 1$  for winning/losing a game, else zero.
- ▶ Atari games  
 $R_t$  implicitly defined as the score in each game

## Agent and Environment




Figure: Slide from David Silver [2015].



## History and state

Define the **history** as the sequence of observations, actions, rewards:

$$H_t = \{O_1, A_1, R_1, \dots, A_{t-1}, R_{t-1}, O_t\}.$$

- ▶ Define the **state**  $S_t$  at time step  $t$  as any function  $f$  of the history:

$$S_t = f(H_t).$$

## History and state

Define the **history** as the sequence of observations, actions, rewards:

$$H_t = \{O_1, A_1, R_1, \dots, A_{t-1}, R_{t-1}, O_t\}.$$

- ▶ Define the **state**  $S_t$  at time step  $t$  as any function  $f$  of the history:

$$S_t = f(H_t).$$

- ▶ The agent has to pick the action  $A_t$  based on  $S_t$ .

## History and state

Define the **history** as the sequence of observations, actions, rewards:

$$H_t = \{O_1, A_1, R_1, \dots, A_{t-1}, R_{t-1}, O_t\}.$$

- ▶ Define the **state**  $S_t$  at time step  $t$  as any function  $f$  of the history:

$$S_t = f(H_t).$$

- ▶ The agent has to pick the action  $A_t$  based on  $S_t$ .
- ▶ Examples:

$S_t = O_t$ , simply the last observation.

$S_t = f(O_{t-k+1}, A_{t-k+1}, R_{t-k+1} \dots, O_t)$ , the last  $k$  transitions of observations, actions and rewards.

## Markov state

We assume that the state representation  $\mathcal{S}$  is **Markov**:

$$\mathbb{P}(S_{t+1}|S_t) = \mathbb{P}(S_{t+1}|S_1, \dots, S_t).$$

- Conclusion: the next action to pick  $A_t$  **only depends** on  $S_t$ .

## Rat example





Figure: Slide from David Silver [2015].

## Rat example



- What if agent state = last 3 items in sequence?
- What if agent state = counts for lights, bells and levers?
- What if agent state = complete sequence?

Figure: Slide from David Silver [2015].

## How to solve the problem

- ▶ So far, we only presented the RL problem, but not how to solve it. A **policy**  $\pi(a|s)$  is the agent's behavior which we want to learn from **experience** (trial-and-error).

## How to solve the problem

- ▶ So far, we only presented the RL problem, but not how to solve it. A **policy**  $\pi(a|s)$  is the agent's behavior which we want to learn from **experience** (trial-and-error).
- ▶ The policy is **stochastic** if  $\pi(a|s) := \mathbb{P}(A_t = a | S_t = s)$ , i.e. we have a distribution over the possible actions to take.



## How to solve the problem

- ▶ So far, we only presented the RL problem, but not how to solve it. A **policy**  $\pi(a|s)$  is the agent's behavior which we want to learn from **experience** (trial-and-error).
- ▶ The policy is **stochastic** if  $\pi(a|s) := \mathbb{P}(A_t = a | S_t = s)$ , i.e. we have a distribution over the possible actions to take.
- ▶ The policy is **deterministic** if  $\pi : \mathcal{S} \rightarrow \mathcal{A}$ , i.e. it's a mapping from state to action.

## How to solve the problem

- ▶ So far, we only presented the RL problem, but not how to solve it. A **policy**  $\pi(a|s)$  is the agent's behavior which we want to learn from **experience** (trial-and-error).
- ▶ The policy is **stochastic** if  $\pi(a|s) := \mathbb{P}(A_t = a | S_t = s)$ , i.e. we have a distribution over the possible actions to take.
- ▶ The policy is **deterministic** if  $\pi : \mathcal{S} \rightarrow \mathcal{A}$ , i.e. it's a mapping from state to action.
- ▶ Note:  $\pi(a|s)$  depend on the **current state**  $S_t$  (not the history).

## Cumulative discount rewards

The **cumulated future discounted rewards** from time  $t \in \{0, 1, \dots, T\}$  is:

$$G_t := \sum_{k=0}^{T-t} \gamma^k R_{t+k}.$$

where  $T$  is the end of the episode.

- ▶  $\gamma \in [0, 1]$  is the discount rate.
  - If  $\gamma$  is close to 0: agent has a "myopic" view, i.e. only short-term rewards are important.
  - If  $\gamma$  is close to 1: agent has a "long-term" view.

## Cumulative discount rewards

The **cumulated future discounted rewards** from time  $t \in \{0, 1, \dots, T\}$  is:

$$G_t := \sum_{k=0}^{T-t} \gamma^k R_{t+k}.$$

where  $T$  is the end of the episode.

- ▶  $\gamma \in [0, 1]$  is the discount rate.
  - If  $\gamma$  is close to 0: agent has a "myopic" view, i.e. only short-term rewards are important.
  - If  $\gamma$  is close to 1: agent has a "long-term" view.
- ▶ Note:  $T$  is a random variable. Ex: end of a game of chess, end of an Atari 2600 game, etc...


## Value function

Define the **value function** of the state  $s$  under the policy  $\pi$ ,  $v_\pi(s) : \mathcal{S} \rightarrow \mathbb{R}$ , as follows:

$$v_\pi(s) := \mathbb{E}_\pi[G_t | S_t = s] = \mathbb{E}_\pi \left[ \sum_{k=0}^{T-t} \gamma^k R_{t+k} | S_t = s \right].$$

- ▶  $v_\pi(s)$  tells us the **goodness/badness** of being in state  $s$ .
- ▶ Intuitively, we want to pick actions which takes us to the best states, i.e. the states where  $v_\pi(s)$  is **large**.

## Toy example (1)



- Rewards: -1 per time-step
- Actions: N, E, S, W
- States: Agent's location

Figure: Slide from David Silver [2015].

## Toy example (2)





Figure: Arrows represent policy  $\pi(a|s)$ .

## Toy example (3)



■ Numbers represent value  $v_\pi(s)$  of each state  $s$



## Value function - Bellman equation

With simple manipulations and the use of the law of iterated expectations, one can show that:

$$\begin{aligned} v_{\pi}(s) &= \mathbb{E}_{\pi} \left[ \sum_{k=0}^{T-t} \gamma^k R_{t+k} | S_t = s \right] \\ &= \mathbb{E}_{\pi} [R_t + \gamma v_{\pi}(S_{t+1}) | S_t = s] \end{aligned} \quad (1)$$

Equation (1) is called the **Bellman equation** for the value function.

## Action-value function

The action-value function  $Q_\pi(s, a) : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  is defined as follows:

$$\begin{aligned} Q_\pi(s, a) &:= \mathbb{E}_\pi \left[ \sum_{k=0}^{T-t} \gamma^k R_{t+k} \mid S_t = s, A_t = a \right] \\ &= \mathbb{E}_\pi [R_t + \gamma Q_\pi(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a] \quad (2) \end{aligned}$$

where (2) is the **Bellman equation** for the action-value function.

## Optimal $Q(s, a)$

- ▶ The **optimal action-value** function  $Q_*(s, a) : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  is defined as:

$$Q_*(s, a) := \max_{\pi \in \Pi} Q_\pi(s, a), \quad \forall (s, a) \in \mathcal{S} \times \mathcal{A},$$

where  $\Pi$  is the set of possible policies.

## Optimal $Q(s, a)$

- ▶ The **optimal action-value** function  $Q_*(s, a) : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  is defined as:

$$Q_*(s, a) := \max_{\pi \in \Pi} Q_\pi(s, a), \quad \forall (s, a) \in \mathcal{S} \times \mathcal{A},$$

where  $\Pi$  is the set of possible policies.

- ▶ Given  $Q_*(s, a)$ , the **optimal** policy  $\pi_*(a|s)$  is *defined* as:

$$\pi_*(a|s) = \begin{cases} 1, & \text{if } a = \operatorname{argmax}_{a \in \mathcal{A}} Q_*(s, a), \\ 0, & \text{otherwise.} \end{cases}$$

Goal: method to obtain  $Q_*(s, a)$ !

# Goal

- 1) Assumption: agent has **no prior knowledge** of the task and the environment to solve.
- 2) In other words: for each pair  $(s, a)$ ,  $Q(s, a)$  is **unknown** and needs to be estimated.

Goal: method to learn  $Q_*(s, a)$  by simulation in order to learn the optimal behavior  $\pi_*$ .

## $\epsilon$ -greedy policy (1)

Suppose  $|\mathcal{A}| = m$ . The type of policy we will work with are called  $\epsilon$ -greedy policy. Let  $\epsilon \in [0, 1]$ . At each time step:

- ▶ Pick a **random** action with probability  $\epsilon$ , with each action having a probability  $\epsilon/m$  of being chosen.

## $\epsilon$ -greedy policy (1)

Suppose  $|\mathcal{A}| = m$ . The type of policy we will work with are called  $\epsilon$ -greedy policy. Let  $\epsilon \in [0, 1]$ . At each time step:

- ▶ Pick a **random** action with probability  $\epsilon$ , with each action having a probability  $\epsilon/m$  of being chosen.
- ▶ Pick the action that maximizes  $Q(s, a)$  with probability  $1 - \epsilon$ , also called the **greedy action**.

## $\epsilon$ -greedy policy (1)

Suppose  $|\mathcal{A}| = m$ . The type of policy we will work with are called  $\epsilon$ -greedy policy. Let  $\epsilon \in [0, 1]$ . At each time step:

- ▶ Pick a **random** action with probability  $\epsilon$ , with each action having a probability  $\epsilon/m$  of being chosen.
- ▶ Pick the action that maximizes  $Q(s, a)$  with probability  $1 - \epsilon$ , also called the **greedy action**.
- ▶ Mathematically:

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a = \underset{A \in \mathcal{A}}{\operatorname{argmax}} Q(s, A), \\ \epsilon/m, & \text{otherwise.} \end{cases}$$



## $\epsilon$ -greedy policy (2)

Why choose an  $\epsilon$ -greedy policy:

- 1) Simplest idea for ensuring **continual exploration** (all  $m$  actions are tried with non-zero probability as long as  $\epsilon > 0$ ).
- 2) Under some constraints on the  $\epsilon$  decay, an  $\epsilon$ -greedy policy converges to  $\pi_*$  (next slides).

## Q-learning(1)

The famous algorithm called **Q-learning** from Watkins [1992] can be used with an  $\epsilon$ -greedy policy to obtain the optimal policy  $\pi_*$ .

- ▶ You will implement a variation of Q-learning in the context of an optimal trade execution problem!

## Q-learning (2)

- 1) Initialize  $Q(s, a)$  randomly  $\forall (s, a) \in \mathcal{S} \times \mathcal{A}$ . Ex:  $Q(s, a) = 0$ .
- 2) Start at the initial state  $s \in \mathcal{S}$ .
- 3) Repeat until the **end of the episode**:
  - 3.1) Select the action  $a \sim \pi(a|s)$ .
  - 3.2) Take action  $a$ , observe the reward  $r$  and move to the state  $s'$ :

$$\{s, a, r, s'\}$$

- 3.3) Update (with  $\alpha \in [0, 1)$ , the **learning rate**)

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \left[ r + \gamma \max_{a'} Q(s', a') \right].$$

- 4) Go back to step 2).

## Q-learning (3)

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \left[ r + \gamma \max_{a'} Q(s', a') \right].$$

Intuition:

- 1)  $Q(s, a)$  is the current estimate of the action-value function for the state  $s$  and action  $a$ .
- 2) The term  $(1 - \alpha)Q(s, a)$  is a weighted value of our current estimate.
- 3) The term  $\alpha \left[ r + \gamma \max_{a'} Q(s', a') \right]$  comes from the Bellman Optimality equation:

$$Q_{\star}(s, a) = \mathbb{E}_{\pi} \left[ R_t + \gamma \max_{A_{t+1}} Q_{\star}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a \right].$$

## Q-learning (4)

Results:

1. Under **certain constraints** on  $\alpha$  and  $\epsilon$ , Q-learning converges to the optimal action-value function  $Q_*$ .

## Q-learning (4)

Results:

1. Under **certain constraints** on  $\alpha$  and  $\epsilon$ , Q-learning converges to the optimal action-value function  $Q_*$ .
2. Q-learning provides the **optimal policy**:

$$a = \operatorname{argmax}_{A \in \mathcal{A}} Q_*(s, A), \quad \forall s \in \mathcal{S}.$$

## Q-learning: Toy example - black board

1. State space:  $\mathcal{S} = \{0, 1, \dots, 4, 5\}$  (6 possible states). Starting state and ending state:  $S_1 = 3$  and  $S_T = 5$ .

## Q-learning: Toy example - black board

1. State space:  $\mathcal{S} = \{0, 1, \dots, 4, 5\}$  (6 possible states). Starting state and ending state:  $S_1 = 3$  and  $S_T = 5$ .
2. Reward matrix (shortest path problem):

$$R = \begin{bmatrix} -1 & 0 & 0 & 0 & -1 & 0 \\ 0 & 0 & 0 & -1 & 0 & 100 \\ 0 & 0 & -1 & -1 & 0 & 0 \\ 0 & -1 & -1 & 0 & -1 & 0 \\ -1 & 0 & 0 & -1 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

where rewards of zero are **impossible** transitions (see board).



## Q-learning: Toy example - black board

1. State space:  $\mathcal{S} = \{0, 1, \dots, 4, 5\}$  (6 possible states). Starting state and ending state:  $S_1 = 3$  and  $S_T = 5$ .
2. Reward matrix (shortest path problem):

$$R = \begin{bmatrix} -1 & 0 & 0 & 0 & -1 & 0 \\ 0 & 0 & 0 & -1 & 0 & 100 \\ 0 & 0 & -1 & -1 & 0 & 0 \\ 0 & -1 & -1 & 0 & -1 & 0 \\ -1 & 0 & 0 & -1 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

where rewards of zero are **impossible** transitions (see board).

3. Initialize  $Q(s, a) = 0, \forall (s, a) \in \mathcal{S} \times \mathcal{A}$ .

## Q-learning: Toy example - black board

1. State space:  $\mathcal{S} = \{0, 1, \dots, 4, 5\}$  (6 possible states). Starting state and ending state:  $S_1 = 3$  and  $S_T = 5$ .
2. Reward matrix (shortest path problem):

$$R = \begin{bmatrix} -1 & 0 & 0 & 0 & -1 & 0 \\ 0 & 0 & 0 & -1 & 0 & 100 \\ 0 & 0 & -1 & -1 & 0 & 0 \\ 0 & -1 & -1 & 0 & -1 & 0 \\ -1 & 0 & 0 & -1 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

where rewards of zero are **impossible** transitions (see board).

3. Initialize  $Q(s, a) = 0, \forall (s, a) \in \mathcal{S} \times \mathcal{A}$ .
4. Additional hyperparameters: fix  $\epsilon = 0.7$ ,  $\gamma = 1.00$  and  $\alpha = 0.9$ .

## Q-learning: Toy example - first transition

- 1) Start at state  $s = 3$ . Three possible actions:  $a \in \{1, 2, 4\}$ .
- 2) Sample action  $a \sim \pi(a|s = 3)$ . Two possibilities:
  - A) Pick a random action with probability 0.7.
  - B) Pick the greedy action with probability 0.3.
- 3) Suppose the action taken is  $a = 4$ . The transition is the following:

$$\{s, a, r, s'\} = \{3, 4, -1, 4\}$$

- 4) Update  $Q(3, 4)$ <sup>1</sup>:

$$\begin{aligned} Q(3, 4) &= (1 - 0.9)Q(3, 4) + 0.9 \left[ -1 + \max_{a' \in \{0, 3, 5\}} Q(4, a') \right] \\ &= (1 - 0.9) \times 0 + 0.9 [-1 + 0] = -0.9. \end{aligned}$$

---

<sup>1</sup>

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \left[ r + \gamma \max_{a'} Q(s', a') \right].$$

## Q-learning: Toy example - second transition

- 1) Currently in state  $s = 4$ . Three possible actions:  $a \in \{0, 3, 5\}$ .
- 2) Sample action  $a \sim \pi(a|s = 4)$ . Three possibilities:
  - A) Pick a random action with probability 0.7.
  - B) Pick the greedy action with probability 0.3.
- 3) Suppose the action taken is  $a = 5$ . The transition is the following:

$$\{s, a, r, s'\} = \{4, 5, 100, 5\}$$

- 4) Update  $Q(4, 5)$ :

$$\begin{aligned} Q(4, 5) &= (1 - 0.9)Q(4, 5) + 0.9 \left[ 100 + \max_{a'} Q(5, a') \right] \\ &= (1 - 0.9) \times 0 + 0.9 [100 + 0] = 90 \end{aligned}$$

## Q-learning: Toy example - conclusion

1. The transitions of the first episode from the starting state  $s = 3$ :

$$3 \longrightarrow 4 \longrightarrow 5.$$

## Q-learning: Toy example - conclusion

1. The transitions of the first episode from the starting state  $s = 3$ :

$$3 \longrightarrow 4 \longrightarrow 5.$$

2. Made the following updates to the Q-function:

$$Q(3, 4) = -0.9, \quad Q(4, 5) = 90,$$

while the other values of  $Q(s, a)$  are **unchanged** and still zero.

## Q-learning: Toy example - conclusion

1. The transitions of the first episode from the starting state  $s = 3$ :

$$3 \longrightarrow 4 \longrightarrow 5.$$

2. Made the following updates to the Q-function:

$$Q(3, 4) = -0.9, \quad Q(4, 5) = 90,$$

while the other values of  $Q(s, a)$  are **unchanged** and still zero.

3. Next steps: start a new episode of training (i.e. start back at the initial state  $s = 3$  and procede the same way).

## Q-learning: Toy example - conclusion

1. The transitions of the first episode from the starting state  $s = 3$ :

$$3 \longrightarrow 4 \longrightarrow 5.$$

2. Made the following updates to the Q-function:

$$Q(3, 4) = -0.9, \quad Q(4, 5) = 90,$$

while the other values of  $Q(s, a)$  are **unchanged** and still zero.

3. Next steps: start a new episode of training (i.e. start back at the initial state  $s = 3$  and procede the same way).
4. Will we converge to the optimal policy? Q-learning converges to  $\pi_\star$  **only if**  $\epsilon$  and  $\alpha$  eventually reaches zero under certain constraints!



# Large-Scale Reinforcement Learning

Most interesting problems are too large to be solved using a tabular method:

- ▶ Backgammon:  $10^{20}$  states;
- ▶ Game of Go:  $10^{170}$  states;
- ▶ Control a robot: continuous state space;

Q-learning **can't be applied** for large-scale problems.

- ▶ In such cases, function approximators for the Q-function (or value function) are often used.
- ▶ For example, a neural network!

## Additional Ressources

- 1) *Reinforcement Learning: An introduction* by R. Sutton and A. Barto (2017). Free book on their website.
- 2) David Silver's online lectures on YouTube (from 2015). Main author on many breakthroughs in reinforcement learning and an excellent teacher!
- 3) Deep Reinforcement Learning course from UC Berkeley <http://rail.eecs.berkeley.edu/deeprlcourse/> with online lectures (Fall 2018).
- 4) For a bridge between RL and optimal control: Course of Dimitri P. Bertsekas (2019) with online lectures <http://web.mit.edu/dimitrib/www/RLbook.html>.